



# 08

Chapiter

# Reinforcement Learning

## 2I1AE1: Artificial Intelligence

Régis Clouard, ENSICAEN - GREYC

“If people do not believe that mathematics is simple,  
it is only because they do not realize how complicated life is.”  
**John Von Neumann, 1947**

# In this chapter

---

2

- Learning by reinforcement
  - In which we examine how an agent can learn from success and failure, from reward and punishment.
- Plan
  - Learning agent
  - Reinforcement Learning
    - ▶ Q-Learning algorithm
    - ▶ Dilemma Exploration / Exploitation
    - ▶ Utility Function Approximation

# 1. Learning Agent

---

3

- An agent **learns** if its performance at a task improves with experience.
- Why programming agent that learns ?
  - An agent can be in an unknown or complex environment (it has to discover it).
  - Even if the environment is known beforehand, it can change overtime in unpredictable ways.
  - Sometime, we have no idea how to program the performance function for an agent.

# Types of Machine Learning

4

## Supervised Learning

Task Driven

*Learning from labeled data “training examples”.*

- Classification: qualitative
- Regression: quantitative

## Unsupervised Learning

Data Driven

*Learning from unlabeled data looking for patterns and structure.*


- Clustering
- Anomaly detection

## Reinforcement Learning

Reward Driven

*Learning by interaction with the environment and outcomes.*

- Decision Making



## 2. Reinforcement Learning

---

5

- What distinguishes reinforcement learning from other automatic learning paradigms:
  - There is no supervisor who indicates the right solution, but rather a reward signal.
  - The reward can be delayed; it is not necessarily instantaneous.
  - The actions of the agent influence the future data the agent will receive.

# Reinforcement Learning Agent

6

```
function INTELLIGENT-AGENT(percept, goal) returns an action
  static: state, the agent's memory of the world state
  state ← UPDATE-STATE-FROM-PERCEPTS(state, percept)
  IF previous_action != None
    LEARN_FROM_TRANSITION(previous_state, previous_action,
                          state, reward)
  action ← CHOOSE-BEST-ACTION(state)
  state ← UPDATE-POLICY-STATE-FROM-ACTION(state, action)
  previous_action ← action
  previous_state ← state
  return action
```

# General Features

7

- We still have a Markov Decision Process (MDP).
- Let
  - $S$  : a set of finite **states** (including an initial state  $s_0$  and terminal states).
  - $A(s)$  : a set of possible **actions** from state  $s$ .
  - $P(s' | s, a)$  : a **transition** model, where  $a \in A(s)$ .
  - $R(s)$ : the **reward** function (how good is to be in state  $s$ ).
  - The environment is perfectly observable.
- We still are looking for the optimal policy  $\pi^*$  that maximizes the expected sum of the rewards.
  - $U^\pi(s) = R(s) + \gamma \sum_{s' \in S} P(s' | s, \pi(s)) U^\pi(s')$  (where  $\gamma$  is the decay discount factor).
  - $\pi^*(s) = \operatorname{argmax}_\pi \sum_{s' \in S} P(s' | s, a) U^\pi(s')$


- New consideration: solve MDP, ie find an optimal policy  $\pi$ , when
  - $P(s' | s, a)$  is unknown
  - $R(s)$  is unknown
- Trial and get information from percepts after choosing an action
  - Percepts provide information on the current state and the immediate reward.



# 3. Q-learning algorithm

9

- **Q-Learning** learns the action-value function  $Q(s, a)$ .
  - $Q(s, a)$  is the expected sum of the rewards from  $s$  and the execution of  $a$  until the end of the optimal policy.
    - ▶  $Q(s, a) = R(s) + \gamma \sum_{s' \in S} P(s' | s, a) U(s')$
  - The link between  $Q(s, a)$  and  $U(s)$  is that  $U(s) = \max_a Q(s, a)$ .
    - ▶  $Q(s, a) = R(s) + \gamma \sum_{s' \in S} P(s' | s, \pi(s)) \max_{a'} Q(s', a')$
- The policy of the agent
  - $\pi(s) = \arg \max_a Q(s, a)$ .
  - Advantage: for the choice of the action, no need to learn  $P(s' | s, a)$  or  $R(s)$  (hidden in  $Q(s, a)$ )



# Learning with Q-learning

10

- According to the definition of  $Q(s_t, a_t)$ , we have:
  - $Q(s_t, a_t) = [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid s_t, a_t]$
  - $R_{t+1}$  is the reward the agent gains after taking action at time step  $t$ .
- We translate this equation by updating based on a learning rate.
  - When a transition  $s \rightarrow s'$  occurs from state  $s$  to state  $s'$ , we apply the following update to  $Q(s, a)$ :
    - ▶  $Q(s, a) \leftarrow Q(s, a) + \underbrace{\alpha}_{\text{learning rate}} (\text{correction})$
    - ▶  $\text{correction} = \text{difference} = \underbrace{R(s) + \gamma \max_{a'} Q(s', a')}_{\text{Current reward + Maximum expected future reward}} - \underbrace{Q(s, a)}_{\text{Previous value}}$
- Theorem: if each action is executed an infinite number of times in each state, the values of  $Q$  lead to the optimal policy.




# Learning Rate $\alpha$

11

- The learning rate  $\alpha$  determines to what extent newly acquired information overrides old information.
  - $Q(s, a) \leftarrow Q(s, a) + \alpha ( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) )$
  - $Q(s, a) \leftarrow (1 - \alpha) Q(s, a) + \alpha (R(s) + \gamma \max_{a'} Q(s', a'))$
  - A factor  $\alpha = 0$  makes the agent learn nothing (exclusively exploiting prior knowledge).
  - A factor  $\alpha = 1$  makes the agent consider only the most recent information (memoryless: ignoring prior knowledge to explore possibilities).
  - In practice, often a constant learning rate is used, such as  $\alpha = 0.1$ .

# Learning with Q-learning. Example

12



## ■ MDP

- $S = \{s_0, s_1, s_2\}$ ,  $A = \{a_0, a_1, a_2\}$

## ■ Initialization:

- $Q(s_0, a_0) = 0$                        $Q(s_0, a_1) = 0$                        $Q(s_0, a_2) = 0$   
 $Q(s_1, a_0) = 0$                        $Q(s_1, a_1) = 0$                        $Q(s_1, a_2) = 0$   
 $Q(s_2, *) = 1$  known terminal node.


## ■ Given:

- $\alpha = 0.5, \gamma = 0.5$

- Note: initial values can be obtained from experiment (eg, helicopter)

# Learning with Q-learning. Example


13



- Observations:  $(s_0)_{-0.1}$ 
  - Nothing to do (we need a triplet  $(s, a, s')$ )
  - Chosen action  $\pi(s_0) = \arg \max\{ Q(s_0, a_0), Q(s_0, a_1), Q(s_0, a_2) \}$   
 $= \arg \max\{ 0, 0, 0 \}$   
 $= a_1$  (arbitrary, could be  $a_0$  or  $a_2$ )

# Learning with Q-learning. Example


14



- Observations:  $(s_0)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1}$ 
  - Recall:  $Q(s, a) \leftarrow Q(s, a) + \alpha ( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) )$  for transition  $s \rightarrow s'$
  - $Q(s_0, a_1) \leftarrow Q(s_0, a_1) + \alpha ( R(s_0) + \gamma \max\{ Q(s_0, a_0), Q(s_0, a_1), Q(s_0, a_2) \} - Q(s_0, a_1) )$   
 $= 0 + 0.5 ( -0.1 + 0.5 \max\{ 0, 0, 0 \} - 0 )$   
 $= -0.05$
  - Chosen action  $\pi(s_0) = \arg \max\{ Q(s_0, a_0), Q(s_0, a_1), Q(s_0, a_2) \}$   
 $= \arg \max\{ 0, -0.05, 0 \}$   
 $= a_0$  (arbitrary, could be  $a_2$ )  
**(change policy! Previously  $\pi(s_0) = a_1$ )**

# Learning with Q-learning. Example


15



- Observations:  $(s_0)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_0} (s_1)_{-0.1}$ 
  - Recall:  $Q(s, a) \leftarrow Q(s, a) + \alpha ( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) )$
  - $Q(s_0, a_0) \leftarrow Q(s_0, a_0) + \alpha ( R(s_0) + \gamma \max\{ Q(s_1, a_0), Q(s_1, a_1), Q(s_1, a_2) \} - Q(s_0, a_0) )$   
 $= 0 + 0.5 ( -0.1 + 0.5 \max\{ 0, 0, 0 \} - 0 )$   
 $= -0.05$
  - Chosen action  $\pi(s_1)$   $= \arg \max\{ Q(s_1, a_0), Q(s_1, a_1), Q(s_1, a_2) \}$   
 $= \arg \max\{ 0, 0, 0 \}$   
 $= a_1$  (arbitrary, could be  $a_0$  or  $a_2$ )

# Learning with Q-learning. Example

16




- Observations:  $(s_0)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_0} (s_1)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1}$ 
  - Recall:  $Q(s, a) \leftarrow Q(s, a) + \alpha ( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) )$
  - $Q(s_1, a_1) \leftarrow Q(s_1, a_1) + \alpha ( R(s_1) + \gamma \max\{ Q(s_0, a_0), Q(s_0, a_1), Q(s_0, a_2) \} - Q(s_1, a_1) )$   
 $= 0 + 0.5 ( -0.1 + 0.5 \max\{ -0.05, -0.05, 0 \} + 0 )$   
 $= -0.0625$
  - Chosen action  $\pi(s_0) = \arg \max\{ Q(s_0, a_0), Q(s_0, a_1), Q(s_0, a_2) \}$   
 $= \arg \max\{ -0.05, -0.05, 0 \}$   
 $= a_2$  **(change policy! Previously  $\pi(s_0) = a_0$ )**



# Learning with Q-learning. Example


17



- Observations:  $(s_0)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_0} (s_1)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_2} (s_1)_{-0.1}$ 
  - $Q(s_0, a_2) \leftarrow Q(s_0, a_2) + \alpha ( R(s_0) + \gamma \max\{ Q(s_1, a_0), Q(s_1, a_1), Q(s_1, a_2) \} - Q(s_0, a_2) )$   
 $= 0 + 0.5 ( -0.1 + 0.5 \max\{ -0.0625, 0, 0 \} + 0 )$   
 $= -0.065625$
  - Chosen action  $\pi(s_1)$   $= \arg \max\{ Q(s_1, a_0), Q(s_1, a_1), Q(s_1, a_2) \}$   
 $= \arg \max\{ 0, -0.0625, 0 \}$   
 $= a_2$  **(change policy! Previously  $\pi(s_1) = a_1$ )**

# Learning with Q-learning. Example


18



- Observations:  $(s_0)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_0} (s_1)_{-0.1} \xrightarrow{a_1} (s_0)_{-0.1} \xrightarrow{a_2} (s_1)_{-0.1} \xrightarrow{a_2} (s_2)_1$ 
  - **Terminal state:**  $Q(s_2, *) = 1$
  - Recall:  $Q(s, a) \leftarrow Q(s, a) + \alpha ( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) )$
  - $Q(s_1, a_2) \leftarrow Q(s_1, a_2) + \alpha ( R(s_1) + \gamma \max\{ Q(s_2, *) \} - Q(s_1, a_2) )$   
 $= 0 + 0.5 ( -0.1 + 0.5 \max\{ 1 \} + 0 )$   
 $= 0.2$
- Let's go for a new try...

## 4. Dilemma Exploration vs. Exploitation

19



- The current approach is greedy: always uses the best action (*exploitation*)
- Problem:
  - Suppose we could do action  $a_3$  in state  $s_1$ , that leads to state  $s_3$  such as  $R(s_3) = 1000$ .
  - Since  $Q(s_1, a_3) = 0$  at initialization, and current  $Q(s_1, a_2) = 0.2$ , after a try  $Q(s_1, a_2) > Q(s_1, a_3)$ , a greedy approach will never explore  $s_3$ !
- Solution : use random actions (exploration)

# Dilemma Exploration vs. Exploitation

---

20

- Analysis
  - Too much exploitation leads to an agent that relies on suboptimal plans.
  - Too much exploration leads to an agent that wastes his time to learn.
- Finding the optimal balance between exploration and exploitation is an open problem in general.
  - Optimal exploration / exploitation strategies exist only in very simple cases.
  - Only practical heuristics are available.

# Solution 1: $\epsilon$ -Greedy

21

- The simplest strategy
  - Consider parameter  $\epsilon \in [0,1]$
  - Every time step, flip a coin  $\rightarrow p$
  - Exploration: with small probability ( $p \leq \epsilon$ ), act randomly
  - Exploitation: with large probability ( $p > \epsilon$ ), act on current policy
- Problems with random actions?
  - You do eventually explore the space, but keep acting randomly once learning is done.
  - One solution: lower  $\epsilon$  over time.

# Solution 2: Exploration Function

22

- A more adaptive strategy: the exploration function  $f(u, n)$ 
    - This function increases artificially the future rewards of unexplored actions to favor exploration.
    - $Q(s, a) \leftarrow Q(s, a) + \alpha( R(s) + \gamma \max_{a'} f( Q[a', s'], N(s', a') ) - Q[a, s] )$   
where  $N(s, a)$  is the number of times action  $a$  has been chosen in state  $s$
- and  $f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \text{ (exploration)} \\ u & \text{otherwise (exploitation)} \end{cases}$
- $R^+$  : an artificial optimistic estimate of the best possible reward obtainable in any state (*problem dependent*)
  - $N_e$  a fix parameter. Guaranty that action  $a$  will be chosen in  $s$  at least  $N_e$  times during learning.
  - The exploration function should decrease for  $u$  and increase for  $n$ .

# Q-learning Algorithm

23

```
function Q-LEARNING-AGENT(percept) returns an action
input: percept, indicating the current state  $s'$  and reward signal  $r'$ 
persistent:  $Q$ , a table of action values index by state and action
              $Nsa$ , a table of frequencies for state-action pairs
              $s, a, r$  the previous state, action, reward, initially null

IF  $s$  is not null THEN
    increment  $Nsa[s, a]$ 
     $Q[s, a] \leftarrow Q[s, a] + \alpha(f(r + \gamma \max_{a'} Q[s', a'], Nsa[s', a']) - Q[s, a])$ 
     $s, a, r \leftarrow s', \arg \max_{a'} f(Q[s', a'], Nsa[s', a']), r'$ 
return  $a$ 
end
```

$f(u, n)$  is the exploration function.

## 5. Limits of this modeling

---

24




- Problem 1: Limited for real applications
  - **Q(state, action)** values are stored in tables of size  $\text{state} * \text{action}$ .
    - ▶ Too large to hold in memory.
    - ▶ Too large to learn all the Q values. Time to convergence increases rapidly as the space gets larger.



# Limits of this modeling

25

- Problem 2: Generalization
- Example: Pacman
  - Suppose we learn during tries that this states is bad (low  $U(s)$  value).
  - With the described Q-Learning approach, we cannot deduce anything about this state.
  - Or even for this one (*one dot less*).



# Solution: Approximation of Utility Function


26

- Describe a state using a vector of features (cf. adversarial search).
  - $f(s) = \sum_i w_i \cdot f_i(s, a)$
  - Features  $f_i$  are floating point functions that capture the salient properties of states.
- Advantages
  - Represent utility functions for a very large number of states.
  - Generalize to similar states.

# Example of State Abstraction: Pacman

27

- Examples of features:
  - Distance to the nearest ghost
  - Distance to the nearest food.
  - Number of ghosts.
  - $1 / (\text{distance to closest dot})$
  - Is Pacman in a tunnel? (0/1)
  - Is Pacman close to a ghost? (0/1)
  - etc.



# Abstraction using Approximated Linear Functions

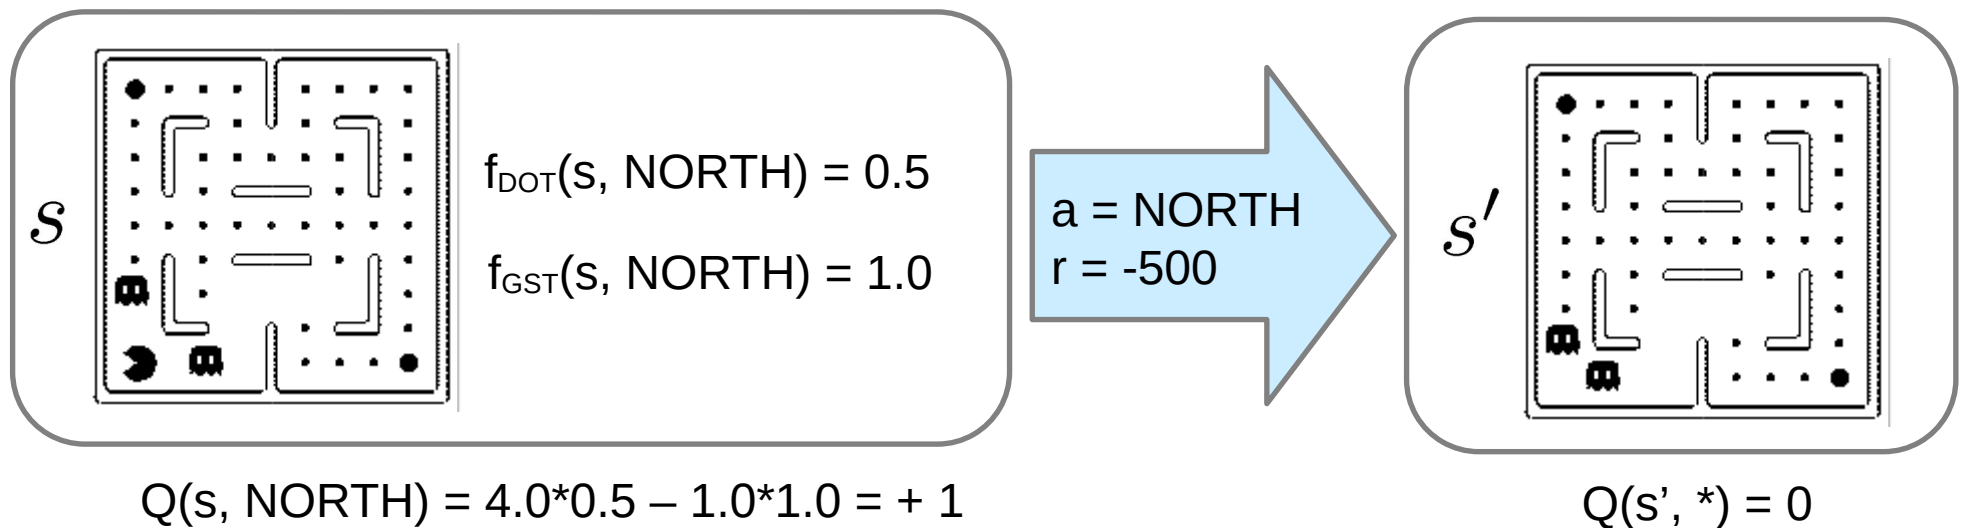
28

- Q values becomes:
  - $Q_f(s, a) = \sum_i w_i \cdot f_i(s, a)$
- Q-Learning learning algorithm is used to learn the weight vector  $W$ , similar to updating Q-values when a transition  $s \rightarrow s'$  by action  $a$  occurs:
  - $\forall i, w_i \leftarrow w_i + \alpha [\text{difference}] \cdot f_i(s, a)$
  - $\text{difference} = [R(s, a) + \gamma \max_{a'} Q(s', a')] - Q(s, a)$
- Intuitive interpretation:
  - Adjust weights of active features
  - e.g., if something unexpectedly bad happens, blame the features that were on: disprefer all states with that state's features.

# Example: Q-Pacman

29

- Suppose 2 features:  $Q(s,a) = 4.0 f_{\text{DOT}}(s, a) - 1.0 f_{\text{GST}}(s, a)$  and  $\alpha = 0.004$ 
  - $f_{\text{DOT}} = 1 / (\text{distance to closest dot})$
  - $f_{\text{GST}} = \text{distance to the nearest ghost}$



- $r + \gamma \max_{a'} Q(s', a') = -500 + 0$
- difference =  $[ R(s, a) + \gamma \max_{a'} Q(s', a') ] - Q(s, a)$   
 $= -500 - 1$

$$w_{\text{DOT}} \leftarrow 4.0 + \alpha[-501] \cdot 0.5 = 3$$

$$w_{\text{GST}} \leftarrow -1.0 + \alpha[-501] \cdot 1.0 = -3$$

- Update:  $Q(s, a) = 3.0 F_{\text{DOT}}(s, a) - 3.0 F_{\text{GST}}(s, a)$

# Conclusion

---

30

- Reinforcement learning is used to learn sequential decision making.
- This is a very active area of research:
  - There are more and more applications in robotics, self-driving, and other areas.
- Reinforcement learning is more difficult when the reward is far (eg. at the very end of a game).
  - Sometimes appropriate positive intermediate reinforcements are added.
    - ▶ Drawback: need expertise.